

# CSE 15L: Software Tools and Techniques Laboratory


Winter 2021 - <http://ieng6.ucsd.edu/~cs15x>

Instructors: Gary Gillespie    Keith Muller

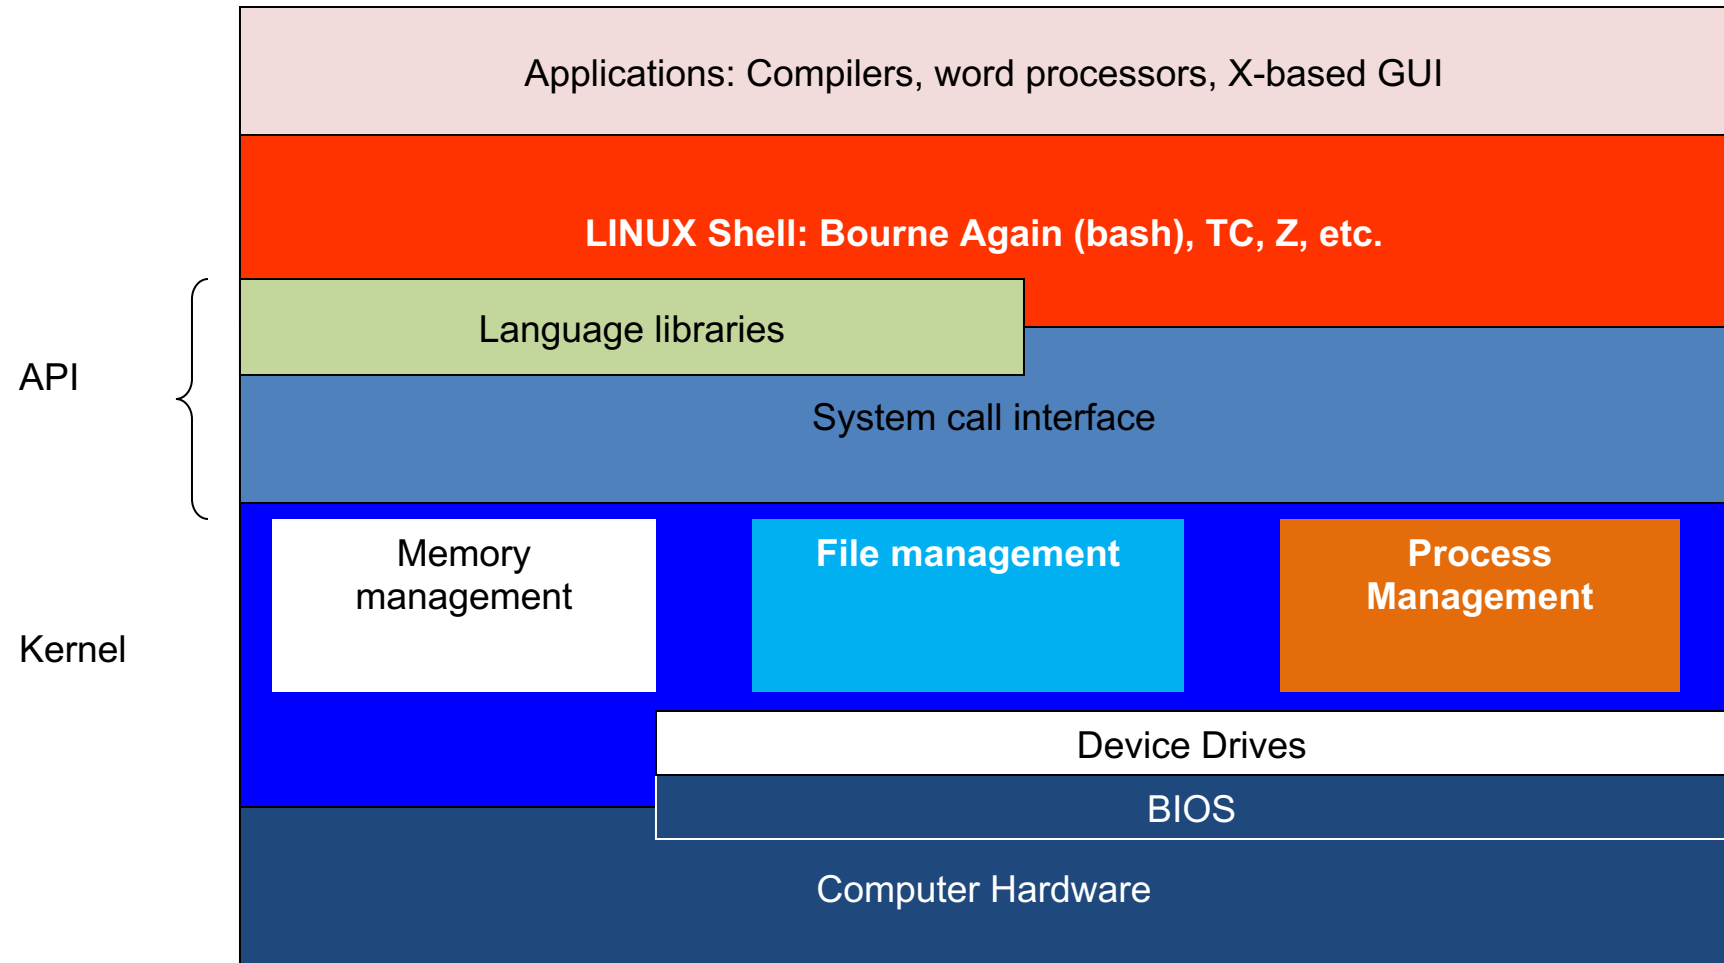
*Class sessions will be recorded and made available to students asynchronously.*

# Operating System

- To facilitate easy, efficient, fair, orderly, and secure use of resources
  - Provide a user interface
  - Organize files on disk
  - Allocating resource to different users with security control
  - Co-ordinate programs to work with devices and other programs



# Abstraction Layers



# Abstraction Layers

- Kernel
  - The part of an OS where the real work is done
- Library and System call interfaces
  - Comprise a set of functions (often known as **A**pplication **P**rogrammer's **I**nterface API) that can be used by the applications and library routines to use the services provided by the kernel
- File Management
  - Control the creation, removal of files and provide directory maintenance
  - For a **multiuser system**, every user should have its own right to access files and directories

# Abstraction Layers

- Memory management software
  - Memory in a computer is divided into **main memory** (RAM) and persistent **secondary storage** (usually refer to disk/flash storage)
- Device driver software
  - Interfaces between the kernel and the BIOS
  - Different device has different driver

# Abstraction Layers

- Process Management


- In a **multitasking system**, multiple programs are executing simultaneously in the system
- When a program starts to execute, it becomes a **process**
- The same program executing at two different times will become two different processes
- The Kernel manages processes in terms of creating, suspending, and terminating them
- A process is protected from other processes and can communicate with the others

# Processes

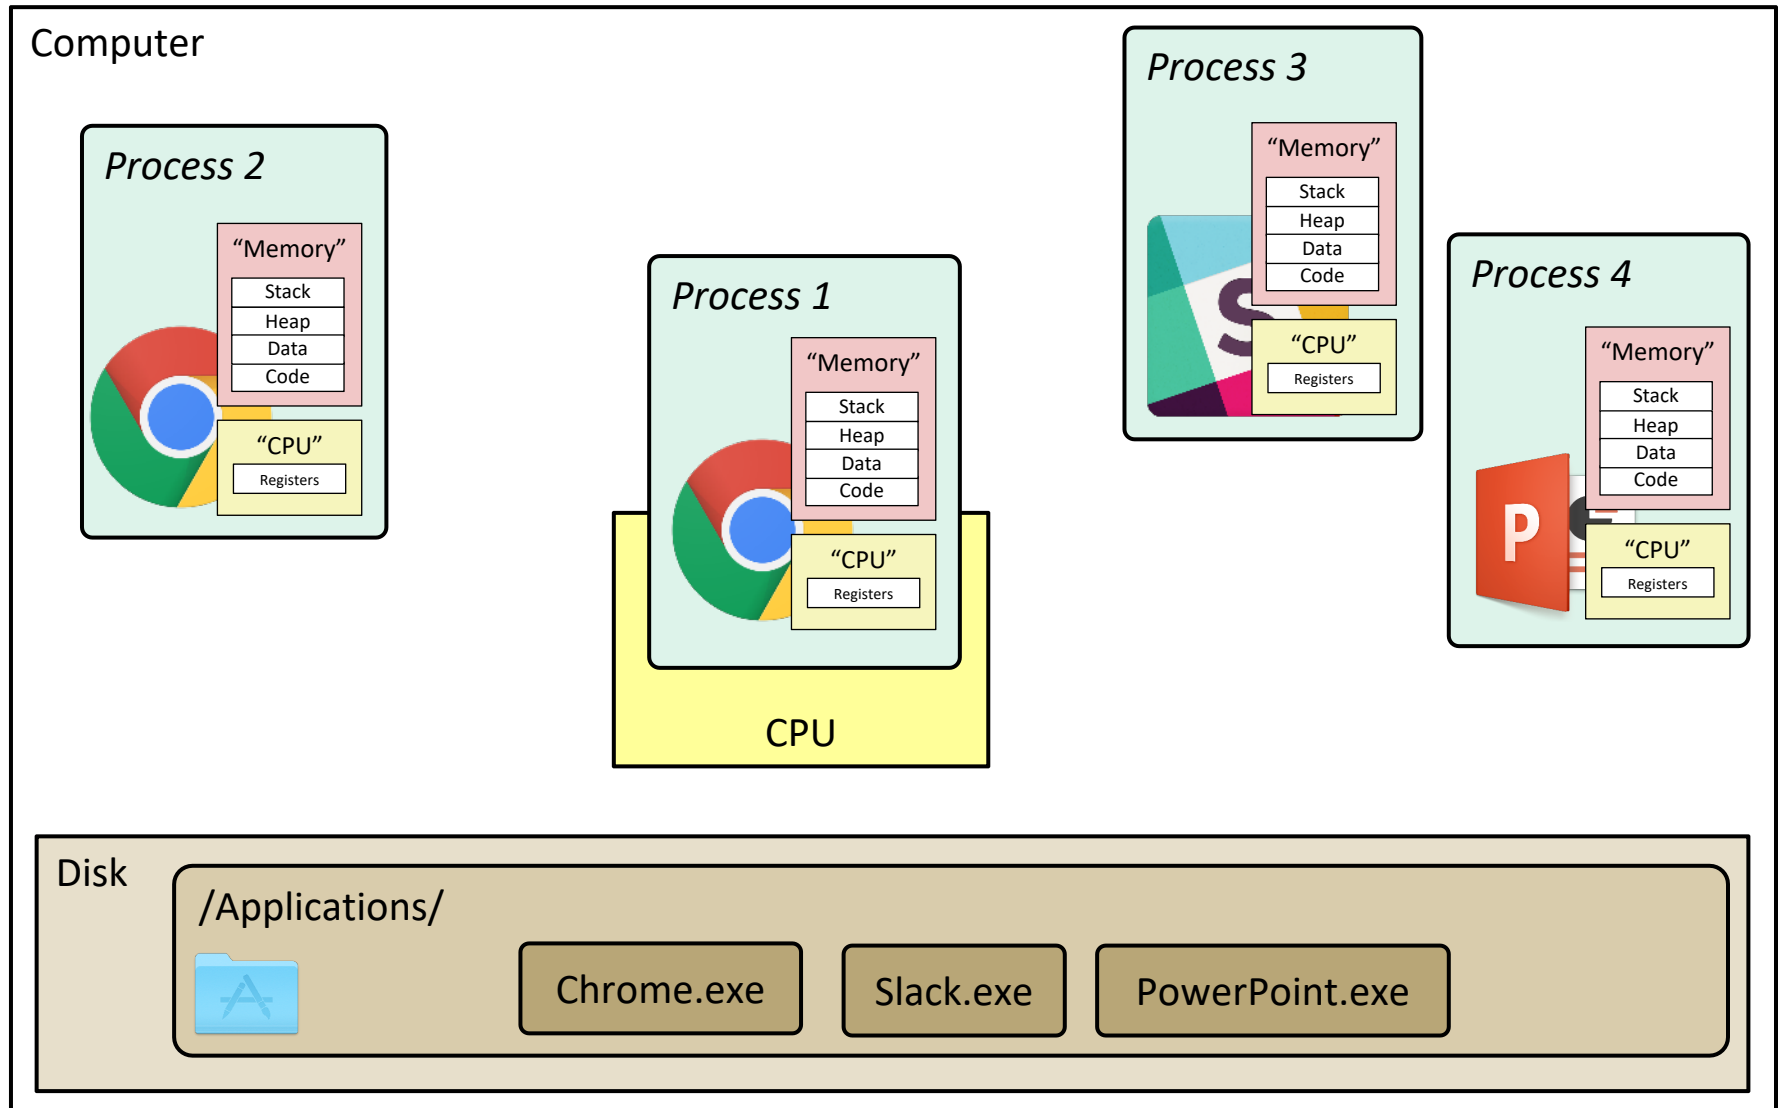
- Informal definition:

A process is a program in execution.

- Process is not the same as a program.
  - Program is a passive entity stored in disk
  - Program (code) is just one part of the process.
- How to start a process?
  - Execute a utility, program, or script!




# What is a process?



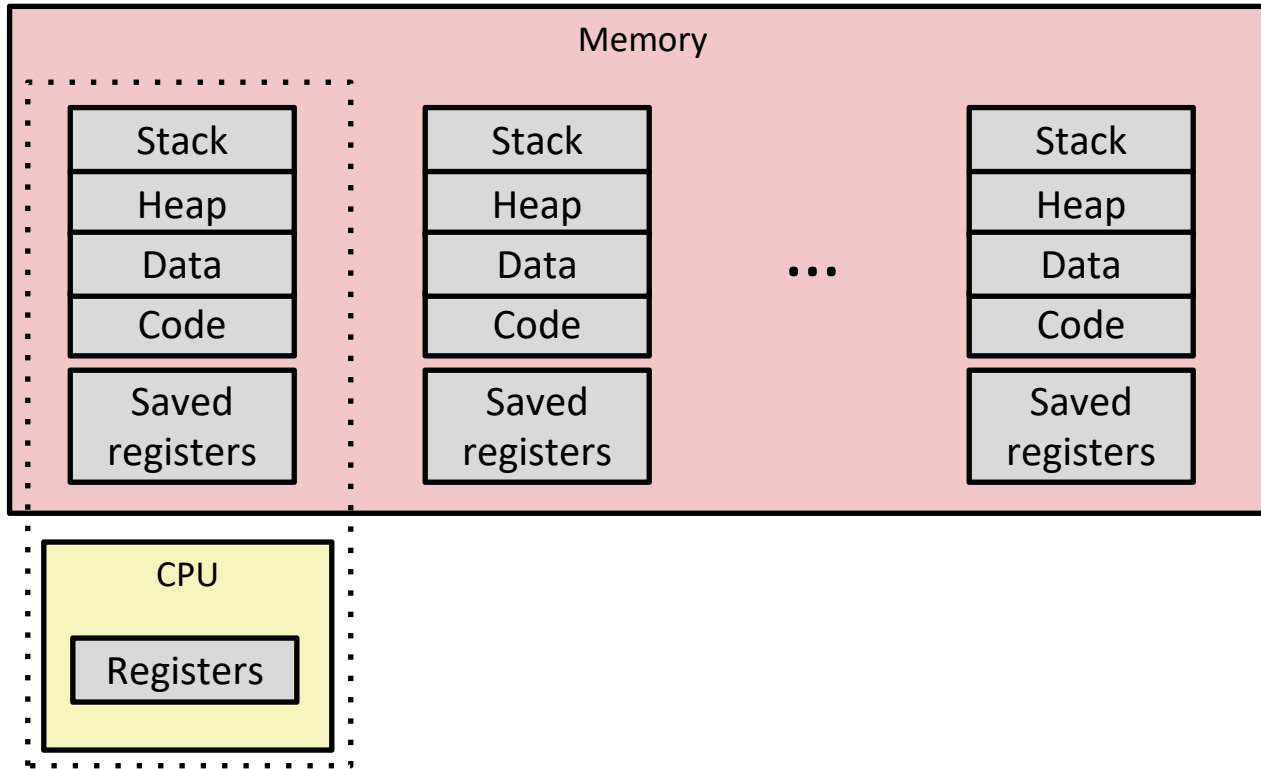


# Multiprocessing: The Illusion



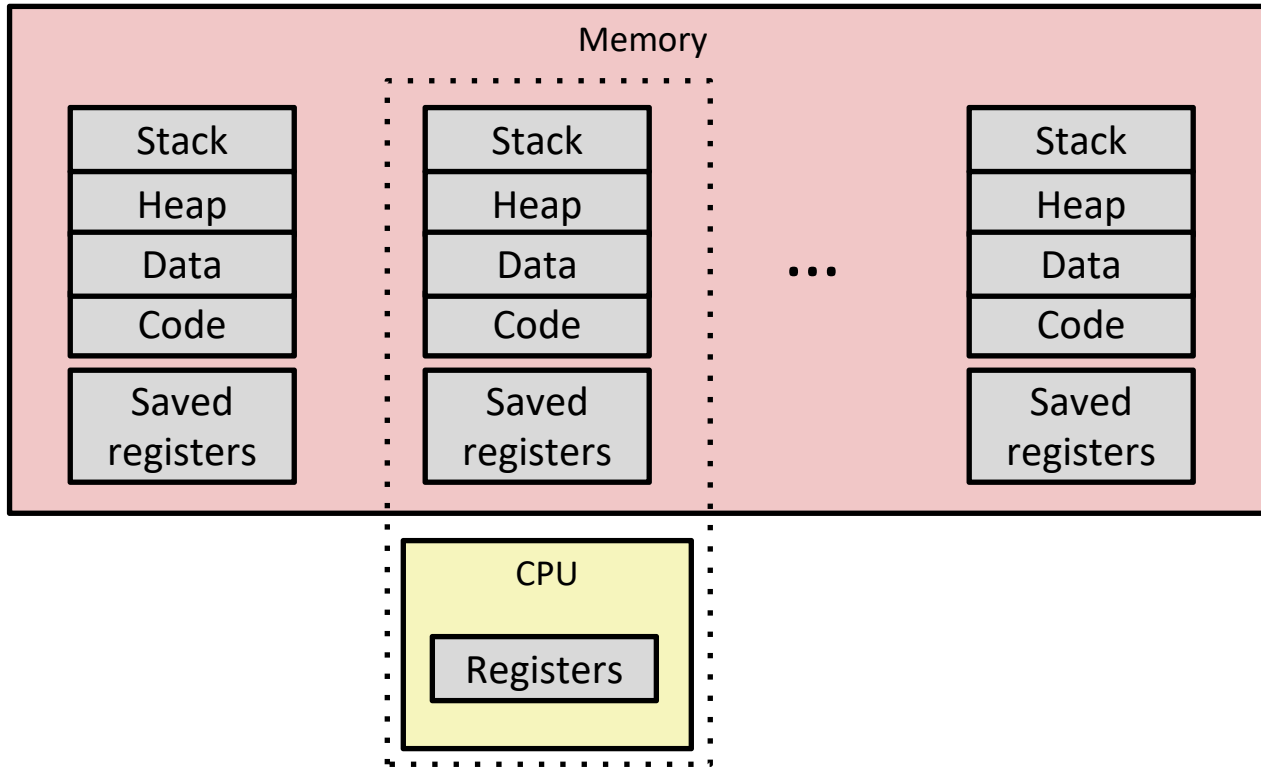
- Computer appears to runs processes concurrently
  - Applications for one or more users
    - Web browsers, email clients, editors, ...
  - Background tasks
    - Monitoring network & I/O devices

# Multiprocessing: The Reality



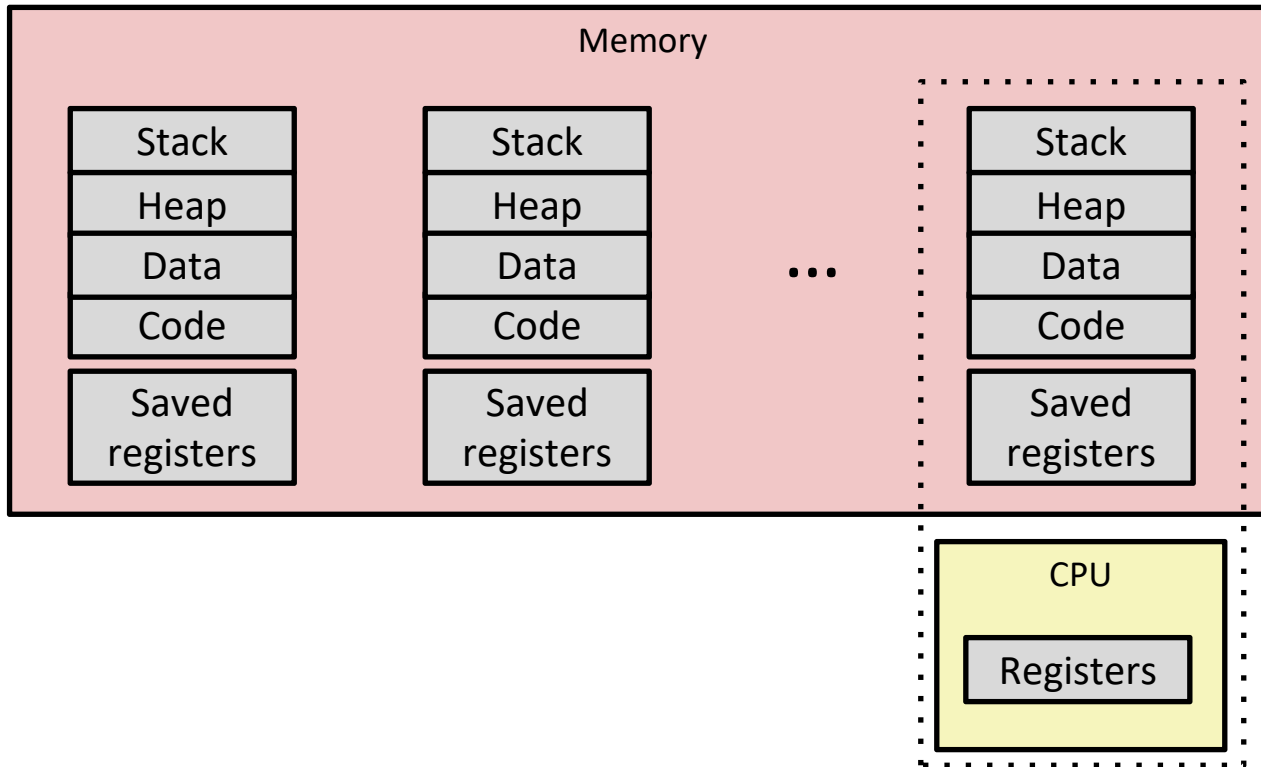
- Single processor executes multiple processes *concurrently*
  - Process executions are interleaved
  - Each CPU runs *one process at a time*

# Multiprocessing: The Reality



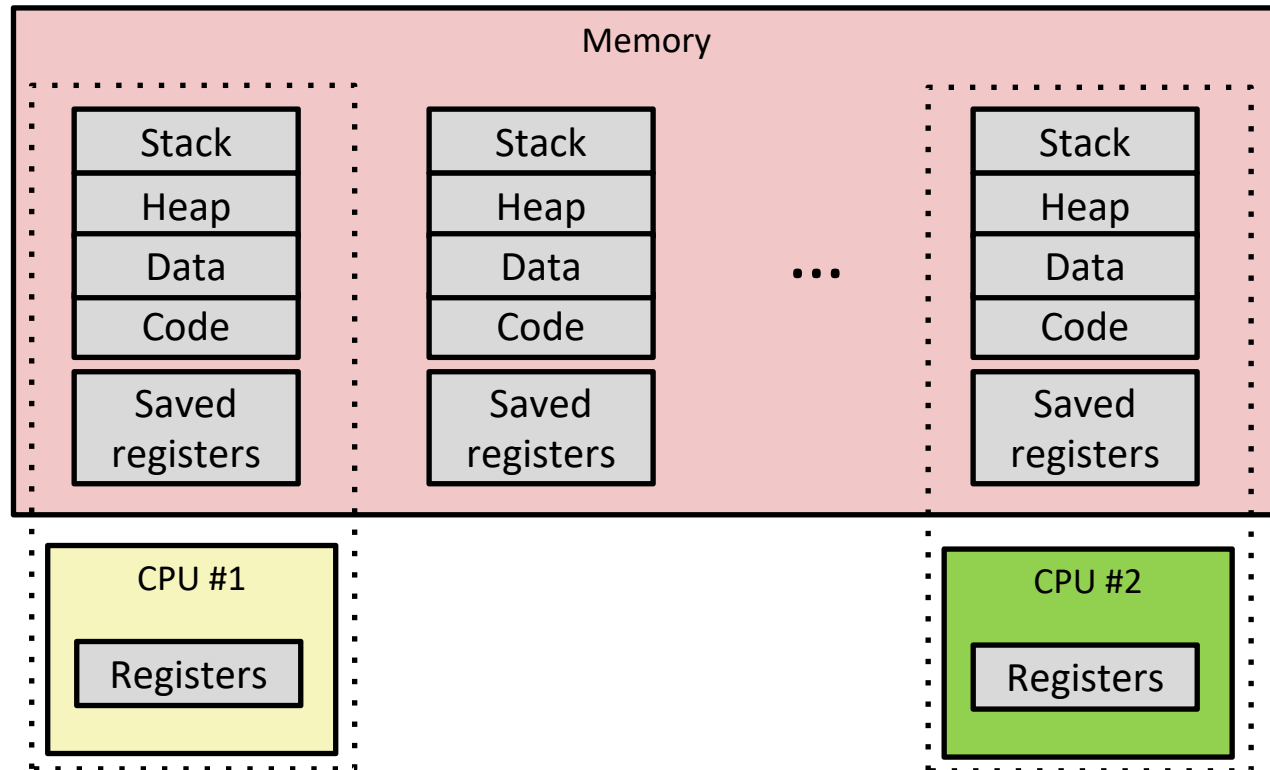
- Single processor executes multiple processes *concurrently*
  - Process executions are interleaved
  - Each CPU runs *one process at a time*

# Multiprocessing: The Reality



- Single processor executes multiple processes *concurrently*
  - Process executions are interleaved
  - Each CPU runs *one process at a time*

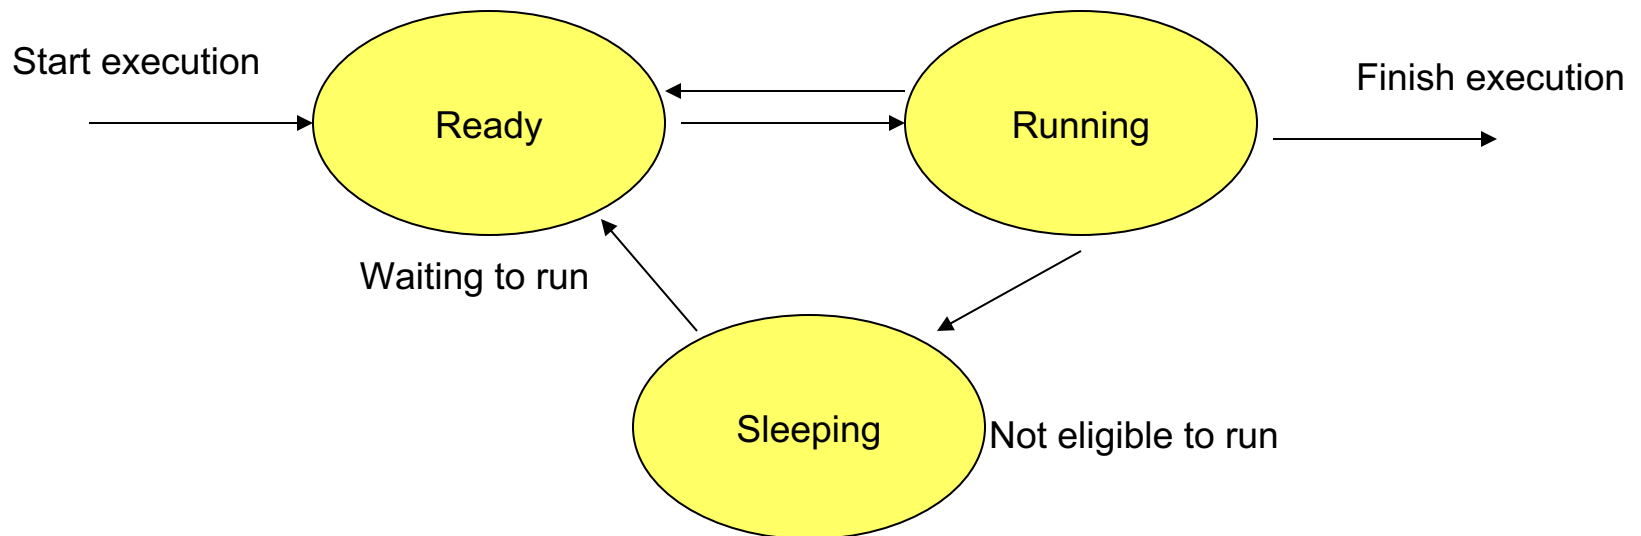
# Multiprocessing: The Reality



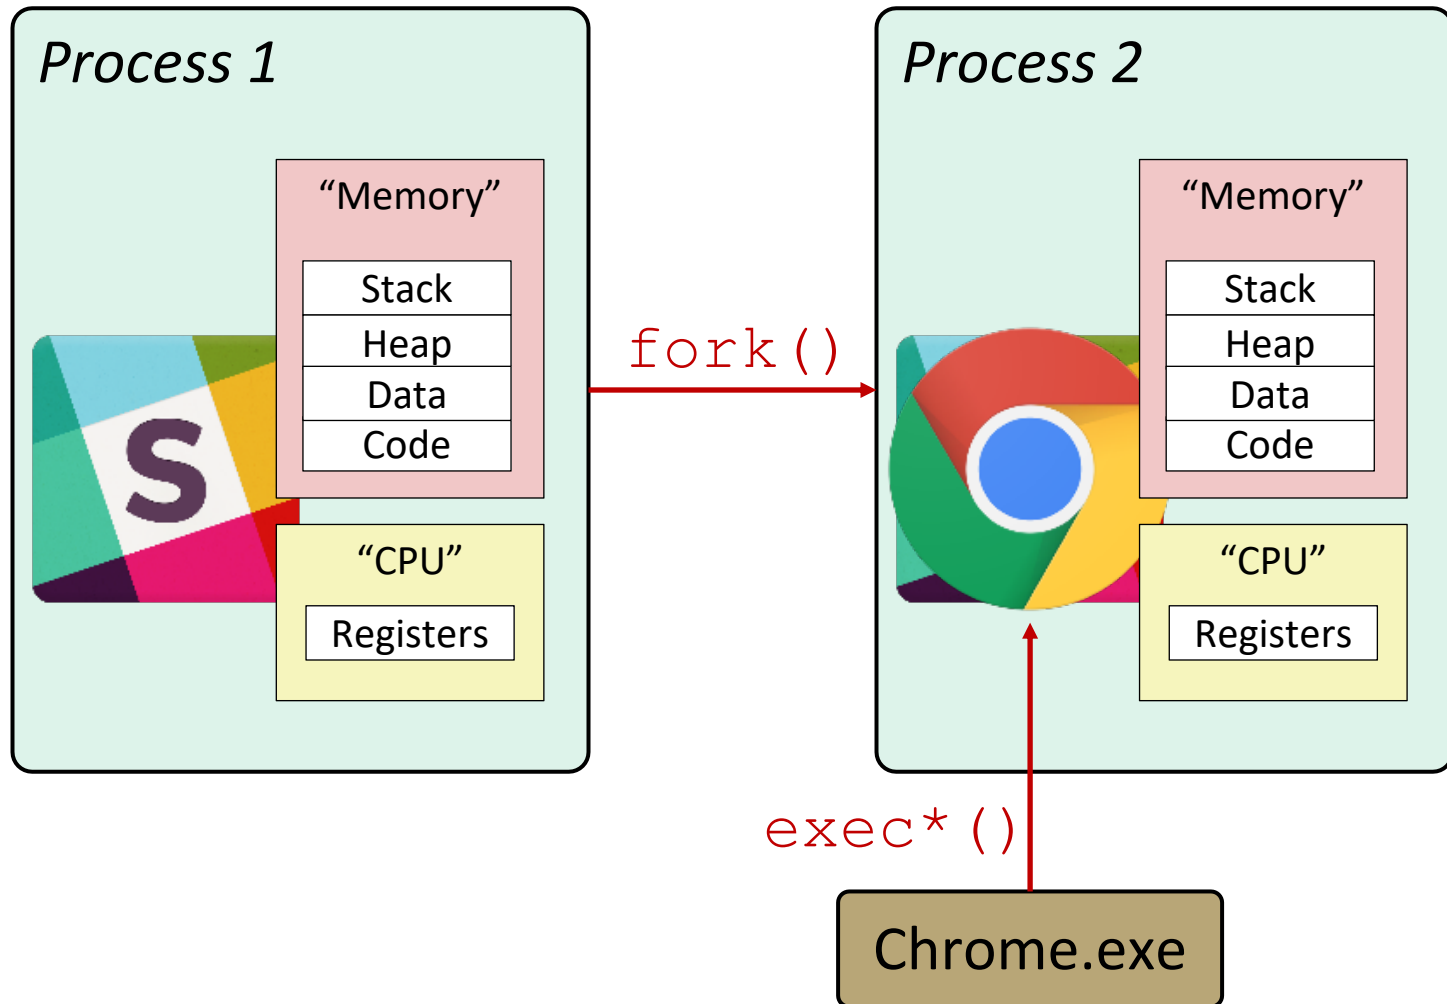
- Multicore processors
  - Multiple CPUs (“cores”) on single chip
  - Shares main memory and I/O devices
  - Each CPU will execute a separate process
    - Kernel schedules processes to cpu cores

# Process States

- A program that is claimed to be executing is called a **process**
- A process is said to be in **at least in one of** the following three **states**:



# Creating New Processes & Programs



# Creating New Processes & Programs

- fork-exec model (Linux):
  - `fork()` creates a copy of the current process
  - `exec*()` replaces the current process' code and address space with the code for a different program
    - Family: `execv`, `execl`, `execve`, `execle`, `execvp`, `execlp`
  - `fork()` and `execve()` are *system calls*
- Parent process has the responsibility to check if the child is done and collect the exit return value
  - `wait()`, `waitpid()`



# The Process ID (PID)

## What is the PID?

- PID: is the Process Identifier Number
- In Linux, uniquely identifies every process
- May be used as input argument to various commands that manipulate processes (**kill**)

# Creating a Child Process

- **Child Process:**

- **fork()** : In Unix, a child process is created when the parent process invoke the **fork()** system call.
- *Child processes are initially identical copies of their parent until they exec() a different program*
- Both parent and child continue execution of the same code after the fork() completes
- *The return value of fork() system call is used by a process to determine if it is the parent or the child*
  - child process – fork returns 0
  - parent process – fork returns PID of child
- When a child process changes “state” (exits, stops execution, etc.), a “**SIGNAL**” is sent to the parent (**SIGCHLD**) the parent can then act

- **Orphan process:**

- When a parent process dies, the child becomes an orphan process.
- Child process are immediately adopted by the **init** process!

# Daemons and Zombies

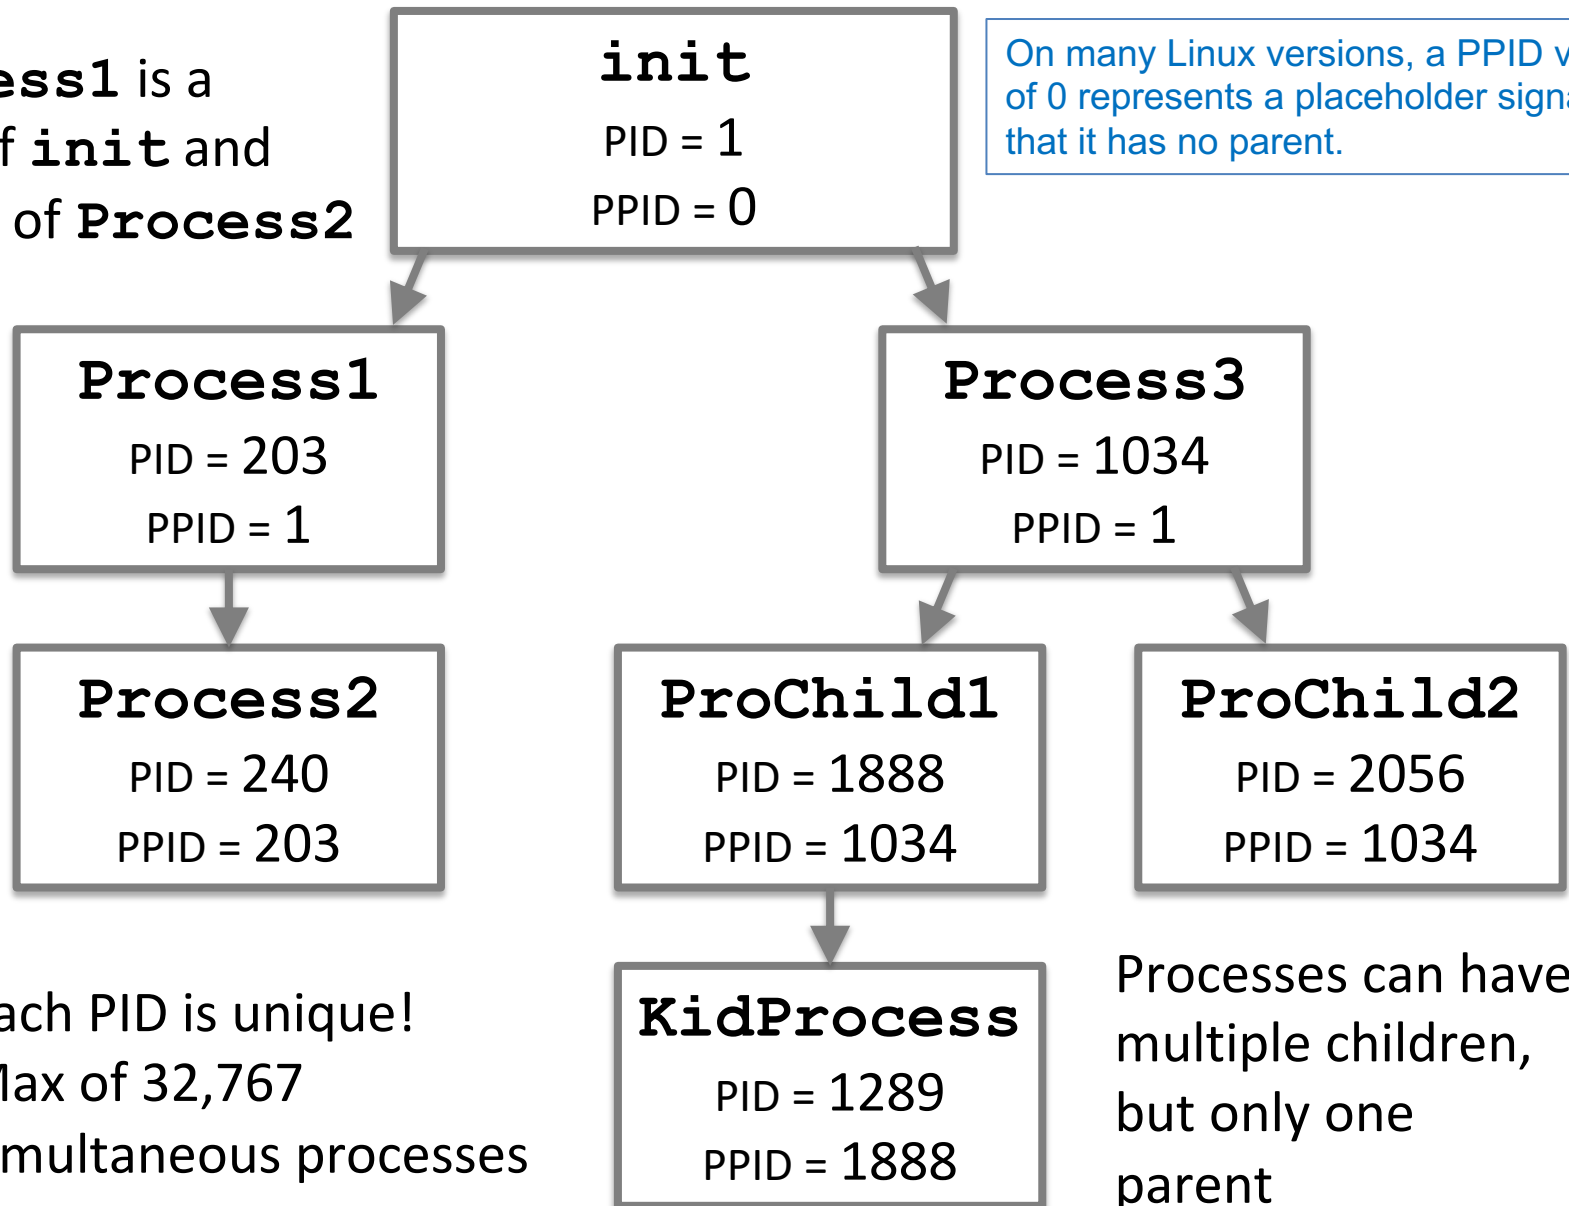
- **Daemon process:**
  - A program run as a background process without being controlled by an interactive user. Sometimes, daemon processes are adopted by `init`.
- **Zombie process**
  - A process that has completed execution, memory and resources are deallocated, but still has an entry in the process table (`z`)
  - Why? it is waiting for the parent to collect its exit status (using the `wait()` or `waitpid()` system call)
  - A zombie is not an orphan!

# Child, parent, and grandparent of all processes...

- Every process has a parent that started it, tracked using PPID (Parent Process ID)
- **init** (Linux) or **launchd** (macOS) is the grandparent of all processes, started when system starts up
  - Always PID 1
- Processes are hierarchical! Think of **init** or **launchd** as the root of the process tree
- **ps tree**: this command shows the relationship of all processes in a tree-like structure

# Example: Process Hierarchy

**Process1** is a child of **init** and parent of **Process2**



On many Linux versions, a PPID value of 0 represents a placeholder signaling that it has no parent.

Each PID is unique!  
Max of 32,767  
simultaneous processes

Processes can have  
multiple children,  
but only one  
parent

# Examining Processes In Linux

- *ps command*
  - Standard process attributes
- */proc directory*
  - More interesting information.
  - Try “man proc”
- *Top, vmstat command*
  - *Examining CPU and memory usage statistics.*

# Simple PS Command

PID	TTY	STAT	TIME	COMMAND
14748	pts/1	S	0:00	-bash
14795	pts/0	S	0:00	-bash
14974	pts/0	S	0:00	vi test1.txt
14876	pts/1	R	0:00	ps ...

Process ID

Terminal name

State:  
S – Sleeping  
(waiting for input)  
R – Running

How much time the process is  
continuously executing

# Example **ps** output

```
$ ps -aef | head
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	12:40	?	00:00:01	/sbin/init splash
root	2	0	0	12:40	?	00:00:00	[kthreadd]
root	4	2	0	12:40	?	00:00:00	[kworker/0:0H]
root	6	2	0	12:40	?	00:00:00	[mm_percpu_wq]

```
$ ps -aef | tail
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
mostert+	2853	2844	0	12:44	pts/0	00:00:00	bash
root	3002	1	0	12:45	?	00:00:00	/usr/sbin/cupsd -l
root	3066	2	0	12:45	?	00:00:00	[kworker/1:0]
root	3077	2	0	12:46	?	00:00:00	[kworker/1:3]
mostert+	3176	2853	0	12:50	pts/0	00:00:00	ps -aef
mostert+	3177	2853	0	12:50	pts/0	00:00:00	tail



# Example **ps** output

```
$ ps -aef | head
```

UID	PID	PPID	C
root	1	0	0
root	2	0	0
root	4	2	0
root	6	2	0

**UID:** User ID that this process belongs to (the person running it)

**PID:** Process ID

**PPID:** Parent process ID (the ID of the process that started it)

**C:** CPU utilization of processor

**STIME:** Process start time

**TTY:** Terminal type associated with the process

**TIME:** CPU time taken by the process

**CMD:** The command that started this process

```
$ ps -aef | tail
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
mostert+	2853	2844	0	12:44	pts/0	00:00:00	bash
root	3002	1	0	12:45	?	00:00:00	/usr/sbin/cupsd -l
root	3066	2	0	12:45	?	00:00:00	[kworker/1:0]
root	3077	2	0	12:46	?	00:00:00	[kworker/1:3]
mostert+	3176	2853	0	12:50	pts/0	00:00:00	ps -aef
mostert+	3177	2853	0	12:50	pts/0	00:00:00	tail

# Shell: Foreground & Background

- **Foreground:**
  - bash forks
  - child bash “execs” into command to execute
  - parent bash sleeps waiting for child to die before prompting
  - Processes executed on terminal run in foreground by default (Receive input from keyboard, Send output to screen)
  - The shell will not run another command until previous command finishes
- **Background:**
  - Invoke a command with an **&** at the end
  - bash forks
  - child bash “execs” into command to execute
  - parent bash doesn’t wait for child to die before prompting for the next command
  - Unless stdin is redirected, a background task will **pause** execution if it tries to read the keyboard
- **nohup: (no hangup)**
  - Used as prefix when starting a background process so you can log off and the process keeps executing

# Listing Current Background Processes (Jobs)

- **jobs**
  - Lists all background processes/jobs in current shell
  - + indicates default for **fg/bg** command

```
$ gnome-calculator &
```

```
[1] 1246
```

```
$ date &
```

```
[2] 1247
```

```
$ Mon Nov 18 14:01:00 PST 2019
```

```
[2]+ Done      date
```

```
$ ./myscript &
```

```
[2] 1269
```

```
$ jobs
```

```
[1]-  Running      gnome-calculator &
```

```
[2]+  Running      myscript &
```

# Managing Foreground/Background

**fg** <#>

- Brings job number <#> to the foreground for keyboard input

**Ctrl-Z**

- Suspends current foreground job temporarily, placing it in the background in a suspended (stopped) condition

**bg** <#>

- Resumes running suspended job <#> in the background

*Optional: Sobell, “A practical Guide to Linux,” pg 150-152, 304-306.*

# Using cntrl-z

```
$ vi a
```

```
(typed a cntrl-z)
```

```
zsh: suspended vi a
```

```
$ vi b
```

```
(typed a cntrl-z)
```

```
zsh: suspended vi b
```

```
$ jobs -l
```

```
[1] - 2817 suspended vi a
```

```
[2] + 2818 suspended vi b
```

```
$ ps 2817
```

PID	TT	STAT	TIME	COMMAND
2817	s000	T	0:00.02	vi a

```
$ ps 2818
```

PID	TT	STAT	TIME	COMMAND
2818	s000	T	0:00.02	vi b

```
$ %1
```

```
(resume execution of vi a)
```

# Example: Background tasks

```
$ (sleep 10; cat > note.txt) &
[1] 985
$ jobs -l
[1]+  985 Running                  ( sleep 10; cat > note.txt ) &
$
[1]+  Stopped                      ( sleep 10; cat > note.txt )
$ fg
( sleep 10; cat > note.txt )
Hi there kitty
$ cat note.txt
Hi there kitty
$
```

Reminder: Ctrl-D indicates end-of-file

# Using NOHUP

```
$ jobs -l
[1]-  1125 Stopped      vi a
[2]+  1126 Stopped      vi b
[3]   1163 Running      ( sleep 10; nohup cat > note.txt ) &
$ nohup: ignoring input and redirecting stderr to stdout
[3]   Exit 1            ( sleep 10; nohup cat > note.txt )
$ cat note.txt
cat: -: Bad file descriptor
$ %+
(resume execution of vi b)
```

# How to kill a process?

## Step 1:

First look at all the processes that are running and filter the ones you want to terminate

```
ps -aef | grep ggillespie
```

This shows the running processes limited to those belonging to **ggillespie**.

## Step 2:

Then kill the process you selected. This command will kill the process with PID number of 125

```
kill 125
```

## Step 3:

Then make sure that the process was killed by looking at the process table again and confirming 125 is not listed

```
ps -aef
```



# kill command

Syntax:

```
kill [-signal] PID
kill [-signal] %JOBID
```

- Can send termination and non-termination signals! But will send **TERM** by default
- Catching a signal allows a process to clean up (removing temp file, etc) before exiting
- If signal isn't trapped, process will die.
- Signal **-9** can't be ignored.
- Only works on processes you own (or when run by **root** user).

# Errors and Signals and Traps

- Consider the following program:

```
echo "this script will endlessly loop until you stop it"
while true; do
    : # Do nothing
done
```

- After you launch this script, it will hang (it is stuck inside a loop!)
- Once started, the script will continue until bash receives a signal to stop it.
- You can send such a signal by typing Ctrl-C which is the signal called **SIGINT** (short for SIGnal INTerupt)

# Common signals...

Signal name	Signal number	Signal description
<b>SIGHUP</b>	1	Hang up detected on controlling terminal or death of controlling process
<b>SIGINT</b>	2	Issued if the user sends an interrupt signal (Ctrl-C)
<b>SIGQUIT</b>	3	Issued if the user sends a quit signal (Ctrl-D)
<b>SIGFPE</b>	8	Issued if an illegal mathematical operation is attempted (division by 0)
<b>SIGKILL</b>	9	If a process gets this signal, it must quit immediately, and will not perform any cleanup operations
<b>SIGALRM</b>	14	Alarm clock signal (used for timers)
<b>SIGTERM</b>	15	Software termination signal (sent by kill by default)

**kill -1** : this command will display all the signals supported by the system.

# trap command

Syntax:

**trap cmd signals**

**signals** is a list of signals to intercept

**cmd** is command to execute when one of signals is received

- **trap** command allows to execute a command when a signal is received by your script

*Optional: Sobell, "A practical Guide to Linux," pg 496-499.*

# Example: **trap** command

```
$ cat talkingloop.sh
```

```
#!/bin/bash
```

```
trap 'echo PROGRAM INTERRUPTED; exit 1' INT
```

```
while true
```

```
do
```

```
    echo "Program running"
```

```
    sleep 1
```

```
done
```

```
$ ./talkingloop.sh
```

```
Program running
```

```
Program running
```

```
Program running
```

```
<Ctrl-C>
```

```
PROGRAM INTERRUPTED
```

*Note: Hard quotes with **trap** cause variables to be evaluated when signal is received. Double quotes cause variables to be evaluated when command is encountered in the script!*

# Next Lecture

- Diagnostic Output
- Java Logging Framework